PROT AI

Project Team

Ibaad Ahmed Chaudhry	22I-0585
Abdullah Kaif Sheikh	22I-2142
Hassan Ali Shoro	22I-0561

Session 2022-2026

Supervised by

Mr.Shoaib Saleem Khattak

Department of Computer Science

National University of Computer and Emerging Sciences
Islamabad, Pakistan

June, 2026

Contents

List of Tables

Chapter 1

Introduction

The purpose of this proposal is to present **ProtAI**, an AI-powered framework for predicting protein dynamics with high accuracy and efficiency. Proteins, often referred to as "locks" in drug discovery, are not static structures—they bend, twist, and reshape themselves as ligands (potential drugs) attempt to bind. Traditional drug discovery methods, relying on static protein snapshots, fail to capture this flexibility, leading to costly and slow pipelines.

ProtAI addresses this challenge by going **beyond static structure prediction**, which tools like AlphaFold have already advanced, to predicting **binding affinity, pocket flexibility, and quantum properties** essential for drug discovery. The system uses **Graph Neural Networks (GNNs)** trained on large-scale datasets such as **MISATO and PDBbind**, refined with quantum mechanics and molecular dynamics. The ultimate aim is to accelerate early-stage drug candidate screening, saving time, money, and lives [**?**].

1.1 Motivation

Drug discovery today is one of the most resource-intensive scientific processes, requiring **10–15 years** and **over $2 billion** to bring a single drug to market, with nearly **90% of candidates failing in clinical trials**. Early-stage failures—often due to inaccurate predictions of protein–ligand binding—waste massive resources.

Recent breakthroughs like **AlphaFold** have solved the challenge of predicting 3D protein structures, proving the power of AI in biology. However, AlphaFold does not address **dynamic binding predictions**, which are crucial for drug discovery. This gap provides a strong opportunity: by combining **AI with Graph Neural Networks**, ProtAI can extend beyond structure prediction to **binding affinity estimation, flexibility prediction, and quantum property modeling**.

The motivation for ProtAI is therefore threefold:

- **Real-world impact** — accelerate drug discovery pipelines and reduce costs.
- **Scientific challenge** — advance the frontier of AI-driven computational biology.
- **Healthcare innovation** — enable faster identification of promising drug candidates, contributing to better and faster medical solutions.

1.2 Problem Statement

The pharmaceutical industry faces a fundamental bottleneck: drug discovery is **too slow, too costly, and highly inefficient**. Current computational methods, such as molecular docking, rely on static protein structures and simplistic scoring functions that fail to capture the **true dynamic and quantum nature** of protein–ligand interactions. Experimental techniques, while accurate, are prohibitively expensive and not scalable.

This creates a critical question: **How can we design a scalable, accurate, and intelligent system for predicting protein–ligand interactions that accounts for protein flexibility and quantum effects?**

ProtAI aims to answer this by developing a system that:

- Predicts **binding affinities** more accurately than docking methods.
- Models **pocket flexibility**, addressing a key limitation of static approaches.
- Integrates **quantum-aware features** such as energy levels and polarizability.
- Scales to large datasets like **MISATO (19,000+ complexes)** for training and generalization.

In short, ProtAI directly tackles the inefficiency and inaccuracy of current methods by introducing an **AI-driven, graph-based approach** capable of bridging the gap between theoretical prediction and practical drug discovery.

1.3 Proposed Solution/Method

ProtAI proposes a **graph neural network–based pipeline** that combines **structural, dynamic, and quantum features** for protein–ligand interaction prediction. The methodology includes:

1.3.1 Data Preparation

- Collect datasets: primarily **MISATO**, along with **PDBbind**.
- Preprocess raw data by cleaning, normalizing, and extracting **atom- and bond-level features**.
- Handle missing values, standardize coordinates, and prepare inputs for graph models.

1.3.2 Graph Representation

- Represent proteins and ligands as graphs: nodes = atoms/residues (with features like charge, atomic number), edges = bonds or spatial interactions.
- Encode 3D distances, bond types, and pocket features.

1.3.3 Graph Neural Network Design

- Implement **message-passing GNNs** (GCNs, GATs) to learn from graph representations.
- Predict key properties:
  - **Binding affinity** (interaction strength).
  - **Binding site flexibility** (structural adaptability).

1.3.4 Validation & Deployment

- Validate on benchmark datasets (e.g., PDBbind).
- Use evaluation metrics such as **Mean Squared Error (MSE)**, **Pearson correlation coefficient (PCC)**, and **AUC**.

This integrated approach goes **beyond AlphaFold** by providing not only static structures but also **interaction dynamics and quantum insights**, making ProtAI a valuable asset for drug discovery pipelines.

Chapter 2

Methodology

2.1 Work Plan and Distribution

The research project is divided into four main phases, spanning a period of 8–9 months. Each phase involves specific tasks and responsibilities assigned to group members. Figure **??** shows the overall timeline, distribution of tasks, and workload allocation.

Phase I: Research & Data Preparation (Month 1–2)

- Conduct literature review, collect MISATO and AlphaFold datasets, and study ML & docking methods.
- Preparation of SRS Document.
- **Work: Group**

Phase II: Feature Engineering & Preprocessing (Month 3–4)

- Extract atom- and bond-level features and construct graph-based datasets.
- Preparation of SRS Document.
- **Work: Hassan and Ibaad**

Phase III: Model Development & Training (Month 5–7)

- Design, train, and validate GNN models for affinity and flexibility prediction.
- Initial Testing Phase.

- **Work: Kaif**

Phase IV: Deployment & Evaluation (Month 8–9)

- Evaluate against benchmarks, and prepare final demo.
- Final Testing Phase.
- **Work: Hassan and Ibaad**

Chapter 3

Preliminary Literature Review

Recent years have witnessed a rapid growth in computational approaches for predicting protein–ligand binding affinity, a key challenge in computer-aided drug discovery. Traditional methods such as molecular docking and scoring functions often struggle with accuracy and scalability due to their reliance on simplified interaction models. To overcome these limitations, researchers have increasingly turned to machine learning and, more recently, deep learning approaches. Among these, graph neural networks (GNNs) and attention-based architectures have shown particular promise, as they can effectively capture the structural and relational complexity of protein–ligand interactions.

This chapter reviews a selection of recent studies that represent different directions in the field, ranging from graph-based frameworks and geometric deep learning to attention-driven architectures and molecular dynamics insights. By examining these works, their strengths, and their limitations, we aim to identify the gaps that motivate the proposed FYP and position it within the broader research landscape.

3.1 Development of a Graph Convolutional Neural Network Model for Efficient Prediction of Protein–Ligand Binding Affinities

Summary: This study presents a Graph Convolutional Neural Network (GCN) for binding affinity prediction. Proteins and ligands are represented as graphs, where nodes correspond to atoms and edges to chemical bonds or interactions. The model leverages convolutional operations to capture local structural dependencies within these graphs.

Critical Evaluation: The key strength of this work lies in its efficient graph representation and scalability compared to traditional docking-based methods. However, its main limitation is the reliance on hand-crafted features for certain interaction descriptors, which may restrict the model's ability to generalize across highly diverse protein families.

Key Takeaway: The paper demonstrates that graph-based deep learning outperforms conventional sequence-based or docking-based predictors in accuracy and efficiency, supporting the case for GNNs in binding affinity prediction.

3.2 Exploring Protein–Ligand Binding Affinity Prediction with Electron Density-Based Geometric Deep Learning

Summary: This work introduces a geometric deep learning framework that incorporates electron density maps of proteins and ligands, enabling the model to capture fine-grained structural and physicochemical information. By treating the 3D electron density as a spatially continuous signal, the model exploits E(3)-invariant operations (rotational and translational invariance) to ensure robust predictions.

Critical Evaluation: The methodology provides a more physics-aware representation of molecular interactions, avoiding the biases introduced by simplified topological graphs. However, the requirement for high-quality electron density data increases computational cost and may limit large-scale application in drug discovery pipelines.

Key Takeaway: This paper highlights the importance of incorporating geometric and physical principles into deep learning for molecular modeling, showing that electron density features can significantly improve predictive performance.

3.3 PLAIG: Protein–Ligand Binding Affinity Prediction Using a Novel Interaction-Based Graph Neural Network Framework

Summary: PLAIG introduces an interaction-focused GNN framework that explicitly models protein–ligand contacts at the interface. Instead of representing protein and ligand separately, it constructs a joint graph capturing inter-molecular edges, allowing the network to directly learn interaction patterns that drive binding affinity.

Critical Evaluation: This approach is innovative because it reduces the need for extensive feature engineering by letting the model learn interaction-specific embeddings. The framework achieves strong benchmark results, but its interpretability is limited, as the black-box nature of GNNs makes it difficult to understand which interaction features are most influential in prediction.

Key Takeaway: PLAIG demonstrates that explicit modeling of protein–ligand interactions can significantly enhance prediction accuracy and may be a promising path for generalizable affinity prediction models.

3.4 Improving Structure-Based Protein–Ligand Affinity Prediction by Graph Representation Learning and Ensemble Learning (2024)

Summary: This study combines **graph representation learning** with **ensemble learning strategies** to improve structure-based binding affinity prediction. Individual GNN models capture local and global structural features, while ensemble methods integrate predictions from multiple learners to enhance robustness and reduce variance.

Critical Evaluation: The hybrid approach improves generalization and reduces overfitting compared to single-model baselines. However, ensemble methods increase computational complexity and training time, which may limit scalability to very large datasets.

Key Takeaway: The work illustrates that combining GNNs with ensemble learning enhances prediction reliability, making it a valuable strategy for real-world drug discovery where data noise and diversity are major concerns.

3.5 Interformer: An Interaction-Aware Model for Protein–Ligand Docking and Affinity Prediction (2024)

Summary: Interformer introduces a **transformer-based model** specifically designed to be interaction-aware. By using attention mechanisms, the model captures long-range dependencies between protein residues and ligand atoms, enabling simultaneous docking pose prediction and binding affinity estimation.

Critical Evaluation: This method effectively addresses the limitation of GNNs in handling long-range interactions. The attention mechanism improves interpretability by highlighting critical binding residues. However, transformer models are computationally expensive, requiring substantial resources for training.

Key Takeaway: The paper demonstrates that transformer architectures, when adapted to protein–ligand systems, provide both improved performance and interpretability in affinity prediction tasks.

3.6 Unveiling Dynamic Hotspots in Protein–Ligand Binding (2023)

Summary: This study applied large-scale molecular dynamics (MD) simulations to identify "dynamic hotspots"—residues and structural regions that strongly influence binding. It analyzed descriptors such as RMSD, solvent-accessible surface area (SASA), and hydrogen-bond occupancies.

Critical Evaluation: The key strength is capturing protein flexibility and time-dependent interactions ignored by static models. The drawback is the high computational cost of MD simulations.

Key Takeaway: Dynamic features provide valuable complementary descriptors for machine learning models in protein–ligand prediction.

3.7 A Brief Review of Protein–Ligand Interaction Prediction (2023)

Summary: This review surveys docking, classical ML, and deep learning methods for protein–ligand interaction prediction. It covers benchmark datasets (PDBbind, MISATO, BindingDB) and evaluation metrics such as RMSE, PCC, and AUC, while discussing dataset bias and generalization challenges.

Critical Evaluation: It provides a strong overview of the state of the art but introduces no new methodology.

Key Takeaway: This work validates the importance of hybrid AI approaches and identifies critical research gaps ProtAI seeks to address.

3.8 DEAttentionDTA: Dynamic Embeddings and Self-Attention for Drug–Target Affinity Prediction (2022)

Summary: DEAttentionDTA introduces dynamic embeddings combined with self-attention to model drug–target affinity (DTA). The attention layers allow the model to capture important residues and atoms contributing to binding.

Critical Evaluation: The model achieves strong performance and interpretability but is computationally demanding and relies on the quality of pretrained embeddings.

Key Takeaway: Self-attention improves both accuracy and interpretability in affinity prediction tasks.

3.9 GraphscoreDTA: Optimized Graph Neural Network for Protein–Ligand Binding Affinity (2023)

Summary: GraphscoreDTA refines GNN architectures for protein–ligand prediction, combining local atomic-level and global structural features. It achieves state-of-the-art results on PDBbind and BindingDB.

Critical Evaluation: Optimized message passing improves predictive power, but the model remains a black box with limited interpretability.

Key Takeaway: This work confirms the central role of optimized GNNs in accurate binding affinity prediction.

3.10 CAPLA: Improved Prediction of Protein–Ligand Binding Affinity by Cross-Attention (2024)

Summary: CAPLA introduces a cross-attention mechanism between protein and ligand embeddings, explicitly modeling inter-molecular interactions.

Critical Evaluation: Cross-attention enables fine-grained interaction modeling, but adds significant computational overhead and requires large training datasets.

Key Takeaway: Cross-attention enhances accuracy in PLI prediction by directly capturing protein–ligand interaction patterns.

3.11 Comparative Tabular Analysis

Paper Title	Research Design	Key Findings	Strengths	Limitations	Relevance to FYP
Development of a GCN Model	Graph convolutional neural network	GCNs outperform docking methods	Efficient, scalable	Relies on hand-crafted features	Establishes baseline for GNN use
Electron Density-Based Geometric DL	Geometric deep learning with E(3)-invariance	Electron density improves accuracy	Physics-aware, rotation-invariant	Requires high-quality data	Highlights physics-based features
PLAIG	Interaction-based GNN	Explicitly modeling interactions improves performance	Learns joint interaction embeddings	Limited interpretability	Supports interaction-focused approaches
Graph Representation + Ensemble (2024)	Hybrid GNN + ensemble	Ensemble improves robustness	Better generalization	Higher computational cost	Shows advantage of ensembles
Interformer (2024)	Transformer-based attention model	Captures long-range dependencies	Interpretability via attention	Computationally expensive	Introduces transformers for PLI prediction
Unveiling Dynamic Hotspots (2023)	MD simulations + analysis	Identifies dynamic residues/features	Captures protein flexibility	Very costly to compute	Adds dynamics to ML feature space
Brief Review (2023)	Survey paper	Summarizes methods, datasets, metrics	Comprehensive overview	No novel method	Frames research gaps for ProtAI
DEAttentionDTA (2022)	Dynamic embeddings + self-attention	Boosts accuracy and interpretability	Attention highlights residues	Resource intensive	Supports use of attention in ProtAI
GraphscoreDTA (2023)	Optimized GNN	Outperforms baseline GCN/GAT	Strong performance, scalable	Limited interpretability	Reinforces GNN optimization
CAPLA (2024)	Cross-attention mechanism	Captures protein–ligand interactions	High accuracy	Computational overhead	Motivates cross-attention in ProtAI

Table 3.1: Comparative analysis of reviewed research papers on protein–ligand binding affinity prediction.

3.12 Summary of Key Insights

From the reviewed literature, several insights emerge:

- **Graph-based approaches are foundational:** GNNs consistently outperform traditional methods.
- **Physics-aware models improve realism:** Incorporating electron density or MD-derived descriptors increases biological fidelity.
- **Attention mechanisms enhance interpretability:** DEAttentionDTA, CAPLA, and Interformer highlight how attention identifies critical residues and improves performance.
- **Ensembles and hybrids improve robustness:** Combining models reduces variance and overfitting.
- **Cross-model integration is the frontier:** Optimized GNNs fused with attention or dynamic features provide the best pathway for generalizable and accurate prediction.

Chapter 4

Conclusions and Future Work

4.1 Conclusions

In this project, we investigated the application of graph-based deep learning techniques for the prediction of protein–ligand binding affinities. The motivation stemmed from the critical role that binding affinity plays in drug discovery and the limitations of traditional computational approaches such as docking and molecular dynamics simulations, which are often time-consuming and computationally expensive.

Through a comprehensive literature review, it was observed that recent advancements in Graph Neural Networks (GNNs), geometric deep learning, and interaction-aware models have significantly improved the ability to capture complex protein–ligand interactions. Models such as GCNs, interaction-based GNNs, and transformer-based architectures highlight the trend towards leveraging structural, geometric, and interaction features for more accurate predictions. These studies not only validated the effectiveness of GNNs but also emphasized the importance of integrating physical and chemical priors into the learning framework.

The outcomes of this work confirm that graph-based models provide a promising pathway for developing efficient, scalable, and interpretable solutions to the binding affinity prediction problem. By modeling proteins and ligands as graphs, and focusing on their interactions, the predictive performance can be substantially improved compared to conventional approaches. This lays a strong foundation for future developments in AI-driven drug discovery pipelines.

4.2 Future Work

While the progress achieved in this domain is encouraging, several avenues remain open for exploration and improvement:

- **Integration of Multi-Modal Data:** Future models could integrate additional biological and chemical information such as electron density maps, protein flexibility, and solvent effects, which are often overlooked in current GNN frameworks.
- **Improved Generalization:** Current models tend to perform well on benchmark datasets but face challenges when applied to novel protein families. Incorporating transfer learning and domain adaptation techniques may help improve robustness and generalizability.
- **Model Interpretability:** Although deep learning models achieve strong performance, their black-box nature limits their usability in drug discovery. Future research should emphasize explainable AI approaches to identify which interactions or features most strongly influence binding affinity predictions.
- **Efficiency and Scalability:** Training large GNNs and transformer-based models can be computationally expensive. Optimization techniques, pruning, and efficient architectures need to be explored for real-world applications involving large-scale datasets.
- **Experimental Validation:** Computational predictions should be coupled with experimental validation in collaboration with wet-lab studies to ensure practical reliability of the proposed models in real drug development pipelines.

In summary, this project establishes the significance of graph-based deep learning for protein–ligand binding affinity prediction. With continued advancements in integrating structural biology, machine learning, and computational efficiency, these methods hold great promise for accelerating drug discovery and precision medicine.

Bibliography

[1] Jinyoung Son and Dongsup Kim. Development of a graph convolutional neural network model for efficient prediction of protein–ligand binding affinities. *Journal of Cheminformatics*, 2021.

[2] Isert Cemens and Li Zhang. Exploring protein–ligand binding affinity prediction with electron density-based geometric deep learning. *Bioinformatics*, 2024.

[3] Hyun Lee and Minsoo Park. Plaig: Protein–ligand binding affinity prediction using a novel interaction-based graph neural network framework. *Briefings in Bioinformatics*, 2023.

[4] Wei Chen and Yong Liu. Improving structure-based protein–ligand affinity prediction by graph representation learning and ensemble learning. *BMC Bioinformatics*, 2024.

[5] Kai Zhao and Mei Huang. Interformer: an interaction-aware model for protein–ligand docking and affinity prediction. *Nature Machine Intelligence*, 2024.

[6] Z. Jin et al. Capla: improved prediction of protein–ligand binding affinity by cross-attention. *Bioinformatics*, 2023.

[7] K. Wang et al. Graphscoredta: optimized graph neural network for protein–ligand binding affinity. *Journal of Cheminformatics*, 2023.

[8] X. Chen et al. Deattentiondta: dynamic embeddings and self-attention for drug–target affinity prediction. *Bioinformatics*, 2024.

[9] A. S. Powers et al. Geometric deep learning for structure-based ligand design. *ACS Central Science*, 2023.

[10] T. Harren et al. Modern machine-learning for binding affinity estimation of structure-based scoring. *Wiley Interdisciplinary Reviews: Computational Molecular Science*, 2024.