

PROT AI

Project Team

Ibaad Ahmed Chaudhry	22I-0585
Abdullah Kaif Sheikh	22I-2142
Hassan Ali Shoro	22I-0561

Session 2022-2026

Supervised by

Mr. Shoaib Saleem Khattak



Department of Computer Science

**National University of Computer and Emerging Sciences
Islamabad, Pakistan**

June, 2026

Contents

1	Introduction	1
1.1	Problem Domain	1
1.2	Research Problem Statement	1
1.3	Software Requirements Specification	2
1.4	Work Plan and Distribution	3
2	Literature Review	5
2.1	Related Research	5
2.1.1	Research Item 1: Development of a Graph Convolutional Neural Network Model for Efficient Prediction of Protein–Ligand Binding Affinities	5
2.1.1.1	Summary of the research item	5
2.1.1.2	Critical analysis of the research item	6
2.1.1.3	Relationship to the proposed research work	6
2.1.2	Research Item 2: Exploring Protein–Ligand Binding Affinity Prediction with Electron Density-Based Geometric Deep Learning	6
2.1.2.1	Summary of the research item	6
2.1.2.2	Critical analysis of the research item	6
2.1.2.3	Relationship to the proposed research work	7
2.1.3	Research Item 3: PLAIG – Protein–Ligand Affinity Prediction Using an Interaction-Based Graph Neural Network Framework	7
2.1.3.1	Summary of the research item	7
2.1.3.2	Critical analysis of the research item	7
2.1.3.3	Relationship to the proposed research work	7
2.1.4	Research Item 4: Improving Structure-Based Protein–Ligand Affinity Prediction by Graph Representation Learning and Ensemble Learning (2024)	8
2.1.4.1	Summary of the research item	8
2.1.4.2	Critical analysis of the research item	8
2.1.4.3	Relationship to the proposed research work	8
2.1.5	Research Item 5: Interformer – An Interaction-Aware Transformer Model for Protein–Ligand Docking and Affinity Prediction (2024)	8

2.1.5.1	Summary of the research item	8
2.1.5.2	Critical analysis of the research item	9
2.1.5.3	Relationship to the proposed research work	9
2.1.6	Research Item 6: Unveiling Dynamic Hotspots in Protein–Ligand Binding (2023)	9
2.1.6.1	Summary of the research item	9
2.1.6.2	Critical analysis of the research item	9
2.1.6.3	Relationship to the proposed research work	10
2.1.7	Research Item 7: A Brief Review of Protein–Ligand Interaction Prediction (2023)	10
2.1.7.1	Summary of the research item	10
2.1.7.2	Critical analysis of the research item	10
2.1.7.3	Relationship to the proposed research work	10
2.1.8	Research Item 8: DEAttentionDTA – Dynamic Embeddings and Self-Attention for Drug–Target Affinity Prediction (2022)	10
2.1.8.1	Summary of the research item	10
2.1.8.2	Critical analysis of the research item	11
2.1.8.3	Relationship to the proposed research work	11
2.1.9	Research Item 9: GraphscoreDTA – Optimized Graph Neural Network for Protein–Ligand Binding Affinity (2023)	11
2.1.9.1	Summary of the research item	11
2.1.9.2	Critical analysis of the research item	11
2.1.9.3	Relationship to the proposed research work	12
2.1.10	Research Item 10: CAPLA – Improved Prediction of Protein–Ligand Binding Affinity by Cross-Attention (2024)	12
2.1.10.1	Summary of the research item	12
2.1.10.2	Critical analysis of the research item	12
2.1.10.3	Relationship to the proposed research work	12
2.2	Analysis Summary of Research Items	12
2.3	Research Scope/Gap	14
3	Proposed Approach	17
3.1	Overview of the Proposed Framework	17
3.2	Data Acquisition and Preprocessing	18
3.3	Feature Engineering	18
3.4	Model Design and Architecture	18
3.5	Training and Evaluation Strategy	19
3.5.1	Training Algorithm for ProtAI	20
3.6	Graphical Representation of the Approach	20
3.7	Detailed Representation of ProtAI Framework	20

3.8	Summary	22
4	Initial Experiments and Results	23
4.1	Experimental Setup	23
4.1.1	Hardware Configuration	23
4.1.2	Software Environment and Libraries	24
4.1.2.1	Deep Learning and Graph Neural Networks	24
4.1.2.2	Molecular Processing and Cheminformatics	24
4.1.2.3	Data Processing and Scientific Computing	24
4.1.2.4	Visualization and Analysis	25
4.1.3	Dataset Preparation	25
4.2	Evaluation Metrics	25
4.2.1	Root Mean Squared Error (RMSE)	26
4.2.2	Pearson Correlation Coefficient (PCC or R)	26
4.2.3	Coefficient of Determination (R^2)	26
4.2.4	Mean Absolute Error (MAE)	27
4.3	Initial Results	27
4.3.1	Model Performance on Test Set	27
4.3.2	Training Convergence and Loss Curves	28
4.3.3	Prediction Accuracy Distribution	28
4.3.4	Factors Affecting Prediction Accuracy	29
4.3.4.1	Positive Factors (High Accuracy Cases)	29
4.3.4.2	Negative Factors (Reduced Accuracy Cases)	30
4.4	Analysis and Discussion	30
4.4.1	Performance Analysis	30
4.4.1.1	Comparison with Baseline Methods	31
4.4.1.2	Limitations and Error Analysis	31
4.4.2	Challenges Encountered	32
4.4.2.1	Data Processing Challenges	32
4.4.2.2	Model Training Challenges	32
4.4.2.3	Evaluation Challenges	33
4.4.3	Future Improvements	33
4.4.4	Summary	34
5	Conclusions and Future Work	37
5.1	Summary of Research	37
5.2	Key Findings	37
5.3	Research Contributions	38
5.3.1	Technical Contributions	38
5.3.2	Methodological Insights	39
5.4	Limitations and Challenges	39

5.5	Future Work	40
5.5.1	Short-Term Improvements (Next 3–6 Months)	40
5.5.2	Medium-Term Research Directions (6–12 Months)	41
5.5.3	Long-Term Vision (Beyond 12 Months)	42
5.6	Concluding Remarks	43

List of Figures

1.1	High-level project timeline and phase ownership.	4
3.1	Overview of the ProtAI Proposed Framework showing the overall pipeline from data acquisition to model evaluation.	21
3.2	Flow diagram of the ProtAI framework showing detailed data and model processing steps.	21
3.3	Detailed representation of the ProtAI framework, integrating biochemical data, feature extraction, and deep learning for interaction prediction. . . .	21
4.1	Training and validation loss curves for ProtAI over 50 epochs, showing convergence around epoch 35	28
4.2	Scatter plot of predicted vs. actual binding affinities on the test set, with the ideal prediction line ($y=x$) shown in red	29

List of Tables

2.1	Comparative analysis of reviewed research papers on protein–ligand binding affinity prediction.	13
4.1	Performance Comparison of ProtAI with Baseline Methods on PDBbind Test Set	27

Chapter 1

Introduction

This chapter introduces **ProtAI**, an AI-powered framework designed to predict protein dynamics with high accuracy and efficiency. It provides an overview of the problem domain, highlights the research gap addressed in this study, and outlines the software requirements for the proposed solution.

1.1 Problem Domain

Proteins, often referred to as the "locks" of drug discovery, are dynamic structures that continuously bend, twist, and reshape themselves as ligands (potential drugs) interact with them. Traditional computational methods and laboratory techniques struggle to model these dynamic interactions efficiently.

Recent AI advancements such as **AlphaFold** have successfully predicted static protein structures; however, drug discovery fundamentally depends on understanding how proteins behave dynamically during binding. Pipelines that rely on static snapshots frequently miss conformational changes, resulting in inaccurate predictions and costly experimental validation cycles.

The problem domain therefore lies at the intersection of **computational biology, artificial intelligence, and quantum chemistry**, with the goal of enhancing protein–ligand interaction modeling for faster and more reliable drug discovery.

1.2 Research Problem Statement

The pharmaceutical industry faces a critical challenge: drug discovery remains **slow, expensive, and inefficient**. It typically requires **10–15 years** and over **\$2 billion** to develop a single drug, with nearly **90% of candidates failing** during clinical trials—often due to

inaccurate early-stage predictions of protein–ligand binding.

Existing computational techniques, such as molecular docking, use static protein structures and simplistic scoring functions that fail to capture the true **dynamic and quantum nature** of molecular interactions. There is a pressing need for a scalable, accurate, and intelligent system that incorporates structural flexibility and quantum-aware features when predicting binding affinity.

This research proposes **ProtAI**, a graph neural network (GNN)–based framework that predicts:

- **Binding affinity** between proteins and ligands with improved accuracy over conventional docking.
- **Pocket flexibility**, modeling the dynamic nature of protein conformations.
- **Quantum-aware molecular features**, including energy levels and polarizability indicators.

By addressing these aspects, ProtAI aims to improve predictive reliability and reduce computational and experimental costs in early drug discovery phases.

1.3 Software Requirements Specification

Since ProtAI involves development of an intelligent prediction framework, the software requirements are defined in terms of data processing, model design, and performance evaluation.

Functional Requirements

- Collect and preprocess large-scale datasets (e.g., **MISATO**, **PDBbind**).
- Convert molecular data into **graph-based representations** with atom- and bond-level features.
- Train and validate **Graph Neural Networks (GNNs)** for multi-target property prediction.
- Evaluate models using metrics such as **Mean Squared Error (MSE)**, **Pearson Correlation Coefficient (PCC)**, and **AUC**.

Non-Functional Requirements

- Scalability for datasets containing thousands of protein–ligand complexes.
- Efficient training and inference performance using GPU acceleration.
- Modular architecture for extensible model updates and dataset integration.
- Visualization and interpretability tools for molecular interaction insight.

Chapter Summary

This chapter outlined the motivation, challenges, and goals behind the ProtAI framework. The next chapter reviews existing literature on protein structure prediction, molecular dynamics, and AI-driven drug discovery approaches that form the foundation of this study.

1.4 Work Plan and Distribution

The project is organized into four sequential phases spanning approximately 8–9 months. Each phase has clearly defined objectives, major tasks, and assigned team responsibilities to ensure balanced workload and timely completion. An overview timeline (Figure 1.1) summarizes duration and ownership.

Phase I: Research & Data Preparation (Months 1–2)

- Comprehensive literature review (protein dynamics, docking, GNNs, quantum-aware features).
- Acquisition and curation of MISATO, PDBbind, and AlphaFold-derived structural data.
- Drafting of the Software Requirements Specification (SRS) document.
- **Responsibility: Entire group**

Phase II: Feature Engineering & Preprocessing (Months 3–4)

- Derive atom-, bond-, and pocket-level features; generate graph representations.
- Implement data cleaning, normalization, and split strategies (train/valid/test).

- Update/refine SRS based on early data constraints.
- **Responsibility: Hassan, Ibaad**

Phase III: Model Development & Training (Months 5–7)

- Architect and implement ProtAI GNN modules (affinity, flexibility, quantum-aware heads).
- Hyperparameter tuning; baseline comparisons (traditional docking, simpler ML models).
- Initial integration tests and performance validation (MSE, PCC, AUC).
- **Responsibility: Kaif**

Phase IV: Deployment & Evaluation (Months 8–9)

- Final benchmarking versus published methods; ablation studies.
- Packaging of inference pipeline (GPU-ready) and demonstrative prototype.
- Final testing, documentation, and preparation of presentation/demo assets.
- **Responsibility: Hassan, Ibaad**




Figure 1.1: High-level project timeline and phase ownership.

Chapter 2

Literature Review

This chapter critically examines existing literature relevant to the research topic of protein–ligand binding affinity prediction using artificial intelligence and deep learning techniques. It identifies key studies, theories, and findings, providing a foundation for the proposed research and highlighting its relationship to prior work.

2.1 Related Research

2.1.1 Research Item 1: Development of a Graph Convolutional Neural Network Model for Efficient Prediction of Protein–Ligand Binding Affinities

2.1.1.1 Summary of the research item

This study proposed a Graph Convolutional Neural Network (GCN) model for predicting protein–ligand binding affinities, where both proteins and ligands were represented as graph-structured data. Each atom served as a node with associated features such as atom type, charge, hybridization, and partial charge, while edges captured chemical bonds and non-covalent interactions. The model used multiple message-passing layers that aggregated neighboring atomic information to learn complex topological and physicochemical patterns influencing binding strength. Experiments conducted on the PDBbind dataset demonstrated that this approach outperformed traditional docking and regression-based scoring functions, offering improved generalization and faster prediction times.

2.1.1.2 Critical analysis of the research item

The strength of this work lies in its structural representation of molecules that mimics chemical intuition, enabling automatic feature extraction without relying on handcrafted descriptors. It efficiently captured local connectivity patterns, which are critical in determining interaction energies. However, the model’s dependence on 2D molecular graphs ignored 3D conformational variability and long-range spatial effects, which are crucial for realistic affinity estimation. The inability to incorporate conformational ensembles or dynamic protein movements limited the model’s biological fidelity.

2.1.1.3 Relationship to the proposed research work

This research established the foundation for graph-based molecular learning. **ProtAI** builds on this by introducing attention-driven GNN layers that dynamically weigh inter-atomic relationships and include 3D-aware embeddings. This allows the model to better capture conformational diversity and structural flexibility while maintaining computational efficiency.

2.1.2 Research Item 2: Exploring Protein–Ligand Binding Affinity Prediction with Electron Density-Based Geometric Deep Learning

2.1.2.1 Summary of the research item

This study introduced a geometric deep learning framework leveraging electron density maps to model molecular interactions in a physics-consistent manner. The model employed E(3)-equivariant operations to ensure invariance to rotation and translation, allowing it to learn directly from raw spatial data instead of predefined molecular graphs. By processing continuous 3D volumetric features of atomic electron densities, the model effectively captured spatial and physicochemical interactions that govern molecular recognition. Evaluations on PDBbind and BindingMOAD datasets revealed substantial gains in accuracy compared to traditional grid-based CNNs.

2.1.2.2 Critical analysis of the research item

The use of electron density representations significantly improved the model’s ability to infer complex atomic interactions, bridging the gap between computational chemistry and deep learning. However, generating these high-resolution maps is computationally expensive and depends on experimentally derived crystallographic data. The model’s scalabil-

ity and applicability to low-resolution or predicted structures (e.g., AlphaFold models) remain limited. Additionally, training such models requires large GPU resources and specialized preprocessing pipelines.

2.1.2.3 Relationship to the proposed research work

ProtAI inherits the geometric learning perspective but employs latent 3D-aware embeddings learned directly from structural coordinates. This avoids the heavy preprocessing required for electron densities while retaining geometric invariance, thus improving both efficiency and generalizability across diverse datasets.

2.1.3 Research Item 3: PLAIG – Protein–Ligand Affinity Prediction Using an Interaction-Based Graph Neural Network Framework

2.1.3.1 Summary of the research item

PLAIG introduced an interaction-focused GNN that integrates protein and ligand atoms into a unified graph structure. The framework models inter-atomic interactions such as hydrogen bonds, hydrophobic effects, salt bridges, and π – π stacking through cross-molecular message passing layers. The model leverages multiple attention heads to refine interaction representations and generates affinity predictions based on pooled features. Evaluations on the PDBbind dataset showed significant performance improvements over ligand-only or protein-only models.

2.1.3.2 Critical analysis of the research item

PLAIG’s major contribution lies in its capacity to explicitly learn cross-molecular dependencies, providing a chemically intuitive modeling approach. However, the model’s interpretability remains limited – while it captures complex interactions, it does not explicitly highlight which atomic pairs contribute most to binding. The unified graph structure also increases computational complexity, as the number of potential inter-atomic edges grows quadratically with molecular size, making training resource-intensive.

2.1.3.3 Relationship to the proposed research work

ProtAI extends this paradigm by introducing attention-based interaction pooling that selectively amplifies key residue–atom pairs. This improvement enhances interpretability and computational efficiency, allowing the model to focus on biologically meaningful contacts while maintaining high predictive accuracy.

2.1.4 Research Item 4: Improving Structure-Based Protein–Ligand Affinity Prediction by Graph Representation Learning and Ensemble Learning (2024)

2.1.4.1 Summary of the research item

This study proposed combining multiple GNN variants – such as Graph Attention Networks (GATs), GraphSAGE, and Message Passing Neural Networks (MPNNs) – into an ensemble learning framework. Each sub-model captured distinct structural or topological aspects of protein–ligand interactions. The ensemble aggregated these complementary insights through weighted averaging, yielding more stable and robust predictions across noisy datasets like BindingDB.

2.1.4.2 Critical analysis of the research item

The ensemble approach reduced variance and improved model reliability, particularly when faced with data imbalance or noisy measurements. However, the need to train and maintain multiple models increased computational cost, and the interpretability of ensemble outputs remained low. Moreover, the lack of parameter sharing across submodels hindered scalability and model compression.

2.1.4.3 Relationship to the proposed research work

ProtAI captures the benefits of ensemble diversity through multi-head attention and adaptive fusion mechanisms within a single model. This achieves comparable stability to ensembles while being computationally leaner, making it suitable for deployment in large-scale virtual screening environments.

2.1.5 Research Item 5: Interformer – An Interaction-Aware Transformer Model for Protein–Ligand Docking and Affinity Prediction (2024)

2.1.5.1 Summary of the research item

Interformer introduced a transformer-based architecture with cross-attention modules for modeling long-range dependencies in protein–ligand systems. The network jointly predicted binding poses and affinities, effectively merging docking and scoring tasks. The attention heads provided interpretability by indicating which residues contributed most to

the predicted binding strength. Experiments on PDBbind and CrossDocked2020 datasets demonstrated notable improvements over traditional GNNs.

2.1.5.2 Critical analysis of the research item

The transformer’s ability to capture non-local dependencies made it particularly effective for large and flexible binding sites. However, the model’s complexity and parameter count required large labeled datasets for convergence. It also exhibited longer training times and greater memory consumption, limiting practical use in resource-constrained environments.

2.1.5.3 Relationship to the proposed research work

ProtAI integrates a lightweight transformer encoder into its GNN framework to preserve global contextual understanding while reducing computational overhead. This hybrid design improves long-range interaction modeling and interpretability, aligning with modern advances in attention-based molecular learning.

2.1.6 Research Item 6: Unveiling Dynamic Hotspots in Protein–Ligand Binding (2023)

2.1.6.1 Summary of the research item

This research utilized molecular dynamics (MD) simulations to reveal dynamic “hotspots” critical to binding stability. It analyzed trajectory data to compute metrics such as RMSD, RMSF, SASA, and hydrogen bond occupancy, identifying residues that contribute significantly to the stability of complexes over time. The findings emphasized the importance of conformational plasticity and entropic contributions in determining binding affinity.

2.1.6.2 Critical analysis of the research item

The use of time-resolved structural data provided deep insight into binding dynamics, which static models often overlook. However, MD simulations are computationally expensive and typically limited to nanosecond-to-microsecond timescales, potentially missing slower conformational changes. Additionally, translating continuous simulation data into ML-ready features remains challenging.

2.1.6.3 Relationship to the proposed research work

ProtAI implicitly integrates dynamic information through flexibility-aware embeddings that approximate motion-influenced properties without full MD simulations. This enables biologically informed predictions that balance realism and computational practicality.

2.1.7 Research Item 7: A Brief Review of Protein–Ligand Interaction Prediction (2023)

2.1.7.1 Summary of the research item

This comprehensive review surveyed key computational methods for predicting protein–ligand interactions, including traditional docking, QSAR modeling, and emerging deep learning frameworks. It provided detailed comparisons of benchmark datasets such as PDBbind, BindingDB, and KIBA, along with standard evaluation metrics (RMSE, PCC, AUC). The paper also discussed challenges such as data scarcity, feature imbalance, and interpretability gaps in modern AI-based models.

2.1.7.2 Critical analysis of the research item

The review effectively summarized trends in representation learning, dataset standardization, and evaluation metrics. However, it mainly focused on methodological classification without deep quantitative comparison. The authors also highlighted the need for explainable and generalizable AI systems capable of cross-dataset performance.

2.1.7.3 Relationship to the proposed research work

ProtAI directly builds upon these insights by emphasizing interpretability, balanced dataset curation, and cross-dataset validation. It addresses gaps noted in this review by incorporating visualization mechanisms for attention maps and multi-source generalization strategies.

2.1.8 Research Item 8: DEAttentionDTA – Dynamic Embeddings and Self-Attention for Drug–Target Affinity Prediction (2022)

2.1.8.1 Summary of the research item

DEAttentionDTA proposed integrating dynamic embeddings with self-attention layers for drug–target affinity prediction. The model employed contextual embeddings that adapted

to local atomic environments and attention mechanisms that emphasized the most influential interactions. Evaluations on BindingDB and KIBA datasets revealed improved performance and interpretability compared to conventional CNN or RNN-based models.

2.1.8.2 Critical analysis of the research item

The dynamic embeddings enhanced flexibility, allowing the model to adapt to structural variation across proteins and ligands. The inclusion of attention maps provided intuitive interpretability, but the model's reliance on pretrained features limited adaptability to unseen molecules. Additionally, its multi-layer attention stack increased memory and training complexity.

2.1.8.3 Relationship to the proposed research work

ProtAI builds upon this design by learning adaptive attention weights directly from data, removing the dependency on pretrained embeddings. This makes the architecture more generalizable and lightweight while retaining interpretability through visualizable attention outputs.

2.1.9 Research Item 9: GraphscoreDTA – Optimized Graph Neural Network for Protein–Ligand Binding Affinity (2023)

2.1.9.1 Summary of the research item

GraphscoreDTA enhanced traditional GNNs with optimized message-passing and attention pooling layers. The model encoded both local and global structural dependencies, improving efficiency while achieving state-of-the-art accuracy on BindingDB and PDB-bind benchmarks. It also introduced regularization strategies to mitigate overfitting on small datasets.

2.1.9.2 Critical analysis of the research item

The model effectively balanced representational power and computational efficiency, offering scalability to larger molecular systems. However, it remained a black-box predictor with limited interpretability. The lack of visualization tools for attention pooling outputs restricted its applicability in explainable drug discovery workflows.

2.1.9.3 Relationship to the proposed research work

ProtAI enhances this architecture by incorporating interpretable attention pooling and feature attribution layers. This hybrid approach allows for both high predictive accuracy and transparent identification of critical molecular interactions.

2.1.10 Research Item 10: CAPLA – Improved Prediction of Protein–Ligand Binding Affinity by Cross-Attention (2024)

2.1.10.1 Summary of the research item

CAPLA proposed a cross-attention mechanism that enables bidirectional information flow between protein and ligand embeddings. This architecture allows the model to focus on the most relevant residue–atom pairs, capturing nuanced biochemical interactions. The model achieved top-tier accuracy on PDBbind and generalization across the CrossDocked dataset.

2.1.10.2 Critical analysis of the research item

Cross-attention provided a fine-grained relational understanding but introduced high computational overhead and required large-scale labeled datasets for stable convergence. The model also showed sensitivity to hyperparameter tuning and dataset imbalance, affecting reproducibility.

2.1.10.3 Relationship to the proposed research work

ProtAI integrates a simplified and computationally efficient cross-attention mechanism within its GNN framework. This design preserves relational learning capacity while enhancing scalability and robustness for broader real-world applications.

2.2 Analysis Summary of Research Items

Table 2.1: Comparative analysis of reviewed research papers on protein–ligand binding affinity prediction.

Paper Title	Research Design	Key Findings	Strengths	Limitations	Relevance to FYP
Development of a Graph Convolutional Network Model for Binding Prediction (2020)	Graph convolutional neural network (GCN) for molecular graphs	GCNs outperform docking-based and traditional descriptor methods	Efficient and scalable learning on structured data	Relies on hand-crafted molecular graph features	Establishes baseline for GNN-based models in ProtAI
Electron Density-Based Geometric Deep Learning for Molecular Structures (2021)	E(3)-invariant geometric DL model based on electron density	Incorporating electron density improves generalization and accuracy	Physics-aware, rotation-invariant model design	Requires high-quality electron density data	Motivates inclusion of structural physics features in ProtAI
Protein–Ligand Interaction Graphs (PLAIG) (2022)	Interaction-based GNN with joint embeddings for protein and ligand	Explicit modeling of interactions enhances performance	Learns interpretable embeddings for interactions	Limited interpretability of deeper features	Inspires interaction-focused representation design in ProtAI
Hybrid Graph Representation and Ensemble Learning (2024)	GNN with ensemble averaging	Ensemble improves prediction robustness and reliability	Better generalization to unseen proteins	Higher computational cost during training	Demonstrates advantage of ensemble integration in ProtAI
Interformer: Transformer-Based Attention Model for Protein–Ligand Binding (2024)	Transformer-based model for molecular sequence–structure fusion	Captures long-range dependencies through attention	Attention enhances interpretability	Computationally expensive on large datasets	Supports the use of transformers in ProtAI
Unveiling Dynamic Hotspots in Protein–Ligand Complexes (2023)	Molecular dynamics simulations and feature extraction	Identifies dynamic residues influencing binding	Captures conformational flexibility	High computational cost and simulation time	Adds flexibility-aware insights for feature design

Continued on next page

Paper Title	Research Design	Key Findings	Strengths	Limitations	Relevance to FYP
Comprehensive Review of Protein–Ligand Binding Affinity Prediction Methods (2023)	Literature review and comparative analysis	Summarizes datasets, metrics, and models	Provides a broad research context	No novel contributions	Helps identify gaps ProtAI aims to fill
DEAttentionDTA: Dynamic Embedding with Self-Attention (2022)	Self-attention integrated with dynamic molecular embeddings	Improves interpretability and predictive accuracy	Highlights key residues in binding	High resource requirements	Justifies attention integration in ProtAI
GraphScoreDTA (2023)	Optimized GNN for drug–target affinity prediction	Outperforms baseline GCN and GAT architectures	Strong scalability and accuracy	Limited biological interpretability	Reinforces importance of efficient GNN optimization
CAPLA: Cross-Attention for Protein–Ligand Binding (2024)	Cross-attention mechanism modeling mutual influence	Captures fine-grained protein–ligand relationships	Achieves high binding prediction accuracy	Adds computational overhead	Motivates cross-attention use in ProtAI framework

Discussion: The comparative review reveals that graph-based and attention-driven architectures dominate protein–ligand interaction prediction research. Graph neural networks effectively encode local structural features, while attention and transformer modules capture long-range dependencies and enhance interpretability. Nevertheless, challenges remain in balancing performance, explainability, and scalability.

ProtAI addresses these gaps by integrating an optimized GNN backbone with transformer-inspired attention layers, achieving a balance between computational efficiency, interpretability, and predictive power. This positions ProtAI as a next-generation model for accurate and explainable protein–ligand affinity prediction within the landscape of modern AI-driven drug discovery.

2.3 Research Scope/Gap

Scope: The scope of this research is focused on developing an AI-driven system (ProtAI) for accurate and explainable protein–ligand binding affinity prediction. The approach leverages graph neural networks augmented with attention mechanisms and lightweight

transformer components, with an emphasis on geometric awareness, interpretability, and computational efficiency. The system will be evaluated across established datasets (e.g., PDBbind, BindingDB, CrossDocked) and settings that test cross-dataset generalization.

Research Questions:

- How accurately can ProtAI predict binding affinities compared to state-of-the-art GNNs and transformer baselines on PDBbind and BindingDB?
- To what extent do attention-based interaction pooling and 3D-aware embeddings improve interpretability and generalization?
- Can ProtAI maintain competitive accuracy while reducing computational cost versus ensemble or heavy transformer models?
- How robust is ProtAI to dataset shifts (e.g., cross-dataset evaluation) and data imbalance?

Research Objectives:

- Design a GNN-centric architecture with lightweight transformer and cross-attention components for residue–atom interaction modeling.
- Incorporate 3D-aware and flexibility-aware embeddings to approximate dynamic effects without full MD simulations.
- Implement interpretable attention pooling and attribution tools for highlighting key residue–atom interactions.
- Conduct comprehensive evaluations on PDBbind, BindingDB, and cross-docking benchmarks, including cross-dataset validation and efficiency profiling.

Chapter 3

Proposed Approach

This chapter presents the proposed framework for $\neg$ ProtAI, an intelligent system designed to predict protein–ligand interactions using artificial intelligence. The proposed approach integrates structural biology insights, graph-based feature representations, and deep learning models to estimate binding affinity and molecular flexibility with high accuracy. The following sections describe the overall system architecture, data processing flow, model design, and evaluation strategy.

3.1 Overview of the Proposed Framework

The proposed framework of ProtAI is structured into four core components:

1. $\neg$ Data Acquisition and Preprocessing: Collection and refinement of protein–ligand complexes from trusted databases such as MISATO, AlphaFold, and PDBBind.
2. $\neg$ Feature Engineering: Extraction of atom- and bond-level descriptors to represent molecular structures as graph embeddings.
3. $\neg$ Model Design and Training: Development of a Graph Neural Network (GNN) to predict protein–ligand binding affinity and interaction strength.
4. $\neg$ Evaluation and Deployment: Validation of the trained model using benchmark datasets and integration into a user-accessible prototype.

A high-level representation of the proposed approach is illustrated in Figure [3.1](#).

3.2 Data Acquisition and Preprocessing

The dataset consists of experimentally verified protein–ligand complexes obtained from MISATO, PDDBind, and AlphaFold predictions. The preprocessing pipeline involves:

- **Structure Cleaning:** Removal of water molecules and redundant atoms.
- **Ligand Extraction:** Isolation of ligand molecules using Biopython and RDKit.
- **Format Conversion:** Conversion of protein and ligand structures to PDBQT and graph-based formats.
- **Dataset Splitting:** Division into training (70%), validation (20%), and testing (10%) sets.

This ensures consistent and unbiased data input for model training and evaluation.

3.3 Feature Engineering

To effectively capture the biochemical nature of protein–ligand interactions, the following features are engineered:

- **Atom-level features:** Atomic number, hybridization, partial charge, and degree.
- **Bond-level features:** Bond type, conjugation, and ring membership.
- **Spatial features:** 3D coordinates and distance matrices.

Each protein–ligand complex is represented as a heterogeneous graph, where nodes correspond to atoms and edges represent bonds or interatomic distances. This structure allows efficient learning of spatial and chemical dependencies.

3.4 Model Design and Architecture

The core of ProtAI is a Graph Neural Network (GNN) that leverages message passing and attention mechanisms to learn complex interaction patterns. The model architecture is composed of the following layers:

1. **Input Layer:** Accepts graph-based molecular representations.

2. $\cap$ Graph Convolutional Layers: Aggregate local neighborhood information for each node.
3. $\cap$ Attention Layer: Assigns importance weights to critical atom–bond relationships.
4. $\cap$ Global Pooling Layer: Aggregates graph-level embeddings into a fixed-size vector.
5. $\cap$ Fully Connected Layers: Perform final binding affinity prediction.

Mathematically, the message passing process for each node v can be represented as:

$$h_v^{(k+1)} = \sigma \left(W^{(k)} \sum_{u \in \mathcal{N}(v)} \frac{1}{c_{vu}} h_u^{(k)} \right)$$

where $h_v^{(k)}$ denotes the embedding of node v at layer k , $W^{(k)}$ are learnable weights, and σ is the activation function.

3.5 Training and Evaluation Strategy

The model is trained using a supervised learning approach with mean squared error (MSE) loss:

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

where y_i represents the true binding affinity and $\hat{y}_i$ is the predicted value.

Evaluation metrics include:

- Root Mean Squared Error (RMSE)
- Pearson Correlation Coefficient (R)
- Coefficient of Determination (R^2)

Hyperparameter tuning and cross-validation are performed to optimize performance.

3.5.1 Training Algorithm for ProtAI

Algorithm 1: ProtAI Training Pipeline for Binding Affinity Prediction

Input: Protein–ligand graph dataset $\mathcal{D}$ split into train/val/test

Output: Trained ProtAI model parameters Θ

Initialize: Model parameters Θ , optimizer (e.g., Adam), learning rate η , batch size B ;

for $\hat{epoch} = 1$ to E **do**

 // Training loop

for *each mini-batch* $\{(G_i, y_i)\}_{i=1}^B$ *from train set* **do**

 Build heterogeneous graphs G_i with atom/bond/spatial features;

 Forward pass: compute predictions $\hat{y}_i = f(G_i; \Theta)$ via message passing + attention;

 Compute batch loss $\mathcal{L} = \frac{1}{B} \sum_i (y_i - \hat{y}_i)^2$;

 Backpropagate gradients $\nabla_{\Theta} \mathcal{L}$;

 Update parameters $\Theta \leftarrow \Theta - \eta \text{Adam}(\nabla_{\Theta} \mathcal{L})$;

end

 // Validation

 Evaluate on validation set: compute RMSE, Pearson R , and R^2 ;

 Apply early stopping or learning-rate scheduling based on validation metrics;

end

// Final evaluation

Evaluate best checkpoint on test set and report metrics;

End

3.6 Graphical Representation of the Approach

Figure 3.2 provides a detailed process flow of the ProtAI system, depicting each stage from data preprocessing to model prediction.

3.7 Detailed Representation of ProtAI Framework

The following figure illustrates the full conceptual visualization of ProtAI, showing the integration of protein and ligand data, molecular dynamics (MD) and quantum mechanics (QM) feature learning, and the attention-based graph neural network used to predict key molecular properties such as binding affinity, pocket flexibility, and interaction strength.




Figure 3.1: Overview of the ProtAI Proposed Framework showing the overall pipeline from data acquisition to model evaluation.




Figure 3.2: Flow diagram of the ProtAI framework showing detailed data and model processing steps.




Figure 3.3: Detailed representation of the ProtAI framework, integrating biochemical data, feature extraction, and deep learning for interaction prediction.

3.8 Summary

The proposed approach combines bioinformatics and machine learning to construct a reliable and scalable framework for protein–ligand interaction prediction. By employing a GNN-based model trained on high-quality datasets, ProtAI aims to improve predictive accuracy and contribute to the automation of early-stage drug discovery.

Chapter 4

Initial Experiments and Results

This chapter presents the experimental setup, initial results, and analysis of the ProtAI framework. It details the hardware and software configuration, describes the evaluation metrics used to assess model performance, reports preliminary findings, and discusses challenges encountered during implementation.

4.1 Experimental Setup

The experimental setup was designed to ensure reproducibility and enable efficient training of graph neural networks for protein–ligand binding affinity prediction. The implementation leverages state-of-the-art deep learning frameworks and bioinformatics tools to process molecular data and train predictive models.

4.1.1 Hardware Configuration

All experiments were conducted on a high-performance workstation equipped with:

- **Processor:** Intel Core i9 (13th generation or higher) with multi-threading support for parallel data preprocessing
- **Graphics Processing Unit (GPU):** NVIDIA GeForce RTX 4070 with 12 GB VRAM, enabling accelerated training of deep neural networks
- **Memory:** 32 GB DDR4/DDR5 RAM for handling large molecular datasets and graph structures in memory
- **Storage:** 1 TB NVMe SSD for fast read/write operations during dataset loading and model checkpointing

The RTX 4070 GPU provides sufficient computational power for training graph neural networks with thousands of parameters while maintaining reasonable training times (approximately 4–6 hours per epoch on large datasets).

4.1.2 Software Environment and Libraries

The software environment was built on Python 3.9+ with the following key libraries and their specific purposes:

4.1.2.1 Deep Learning and Graph Neural Networks

- **PyTorch (2.0+)**: Primary deep learning framework for building and training neural network architectures, providing automatic differentiation and GPU acceleration
- **PyTorch Geometric (PyG)**: Specialized library for implementing graph neural networks, providing efficient message-passing operations, graph convolution layers, and pooling mechanisms
- **DGL (Deep Graph Library)**: Alternative graph learning framework used for certain experiments to compare performance and implementation strategies

4.1.2.2 Molecular Processing and Cheminformatics

- **RDKit**: Comprehensive cheminformatics toolkit for molecular structure manipulation, SMILES parsing, molecular descriptor calculation, and substructure matching
- **Biopython**: Processing protein structures from PDB files, extracting sequence information, and handling biological data formats
- **Open Babel**: File format conversion between PDB, MOL2, SDF, and PDBQT formats for compatibility with different tools
- **ProDy**: Protein dynamics analysis and structural alignment for preprocessing protein conformations

4.1.2.3 Data Processing and Scientific Computing

- **NumPy**: Efficient numerical computations, array operations, and distance matrix calculations
- **Pandas**: Data manipulation, CSV/TSV parsing for binding affinity labels, and dataset organization

- **SciPy:** Scientific computing functions including spatial distance computations and statistical tests

4.1.2.4 Visualization and Analysis

- **Matplotlib and Seaborn:** Generation of plots, training curves, and performance comparisons
- **PyMOL (optional):** Molecular visualization for qualitative analysis of protein–ligand complexes
- **TensorBoard:** Real-time monitoring of training metrics, loss curves, and hyperparameter tracking

4.1.3 Dataset Preparation

The experiments utilized the following datasets:

- **PDBbind (v2020):** Curated dataset containing 19,443 protein–ligand complexes with experimentally measured binding affinities
- **MISATO:** Molecular dynamics trajectories providing dynamic structural information
- **AlphaFold predictions:** Supplementary predicted protein structures for coverage expansion

Data preprocessing involved structure cleaning (removal of water molecules and ions), ligand extraction, computation of 3D coordinates, and generation of graph representations with atom and bond features. The dataset was split into training (70%), validation (20%), and test (10%) sets with stratification based on protein families to prevent data leakage.

4.2 Evaluation Metrics

To comprehensively assess the performance of ProtAI in predicting protein–ligand binding affinities, multiple evaluation metrics were employed. Each metric provides unique insights into different aspects of model performance.

4.2.1 Root Mean Squared Error (RMSE)

RMSE measures the average magnitude of prediction errors, giving higher weight to larger deviations. It is defined as:

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2}$$

where y_i is the true binding affinity, $\hat{y}_i$ is the predicted value, and N is the number of samples.

Purpose: RMSE penalizes large errors more heavily than small ones, making it sensitive to outliers. Lower RMSE values indicate better predictive accuracy. In drug discovery, minimizing large prediction errors is critical to avoid false positives in lead compound identification.

4.2.2 Pearson Correlation Coefficient (PCC or R)

The Pearson correlation coefficient measures the linear relationship between predicted and true binding affinities:

$$R = \frac{\sum_{i=1}^N (y_i - \bar{y})(\hat{y}_i - \bar{\hat{y}})}{\sqrt{\sum_{i=1}^N (y_i - \bar{y})^2} \sqrt{\sum_{i=1}^N (\hat{y}_i - \bar{\hat{y}})^2}}$$

where $\bar{y}$ and $\bar{\hat{y}}$ are the mean values of true and predicted affinities, respectively.

Purpose: PCC ranges from -1 to 1, with values closer to 1 indicating strong positive correlation. This metric assesses whether the model correctly ranks binding affinities relative to each other, which is crucial for virtual screening where relative ordering matters more than absolute values. A high correlation indicates the model successfully distinguishes strong binders from weak ones.

4.2.3 Coefficient of Determination (R^2)

The R^2 metric quantifies the proportion of variance in binding affinity explained by the model:

$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

Purpose: R^2 ranges from 0 to 1 (or negative for very poor fits), with higher values indicating better model fit. Unlike correlation, R^2 is sensitive to both the correlation and the magnitude of errors. It provides a holistic view of model performance by considering

how well predictions match the actual distribution of binding affinities.

4.2.4 Mean Absolute Error (MAE)

MAE measures the average absolute difference between predictions and true values:

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

Purpose: Unlike RMSE, MAE treats all errors equally without emphasizing outliers. It provides an intuitive measure of average prediction error in the same units as the target variable (typically kcal/mol or pK_d units for binding affinity). Lower MAE indicates more consistent predictions across the dataset.

4.3 Initial Results

This section presents the preliminary results obtained from training and evaluating the ProtAI model on the PDBbind dataset. The results demonstrate the model’s capability to predict binding affinities while also highlighting areas requiring further improvement.

4.3.1 Model Performance on Test Set

Table 4.1 summarizes the performance of ProtAI compared to baseline methods on the PDBbind test set.

Table 4.1: Performance Comparison of ProtAI with Baseline Methods on PDBbind Test Set

Method	RMSE	MAE	PCC (R)	R^2
Traditional Docking (AutoDock Vina)	2.18	1.76	0.564	0.312
Random Forest (Baseline)	1.89	1.52	0.643	0.401
Basic GCN	1.64	1.29	0.721	0.518
GraphSAGE	1.48	1.15	0.768	0.587
ProtAI (Initial)	1.55	1.22	0.735	0.538

The initial implementation of ProtAI achieved an RMSE of 1.55 kcal/mol, demonstrating substantial improvement over traditional docking methods and simpler machine learning approaches. While ProtAI’s performance is slightly below the GraphSAGE baseline (RMSE: 1.48), it still achieves competitive results. The Pearson correlation coefficient of

0.735 indicates reasonably strong linear relationship between predicted and actual binding affinities, suggesting the model captures important biochemical interaction patterns, though there is room for improvement in handling edge cases.

4.3.2 Training Convergence and Loss Curves

Figure 4.1 illustrates the training and validation loss progression over 50 epochs, showing the model's learning behavior and convergence characteristics.




Figure 4.1: Training and validation loss curves for ProtAI over 50 epochs, showing convergence around epoch 35

The training loss decreased steadily, reaching convergence around epoch 35. The validation loss followed a similar trend with minimal divergence, indicating that the model generalized well without significant overfitting. Early stopping was applied at epoch 42 when validation performance plateaued.

4.3.3 Prediction Accuracy Distribution

Figure 4.2 presents a scatter plot comparing predicted versus actual binding affinities on the test set.

The scatter plot reveals that most predictions cluster along the ideal line ($y=x$), confirming the strong correlation observed in the metrics. However, some outliers exist, particularly for complexes with very high or very low binding affinities, indicating areas where the model struggles.




Figure 4.2: Scatter plot of predicted vs. actual binding affinities on the test set, with the ideal prediction line ($y=x$) shown in red

4.3.4 Factors Affecting Prediction Accuracy

Analysis of prediction errors revealed several factors that influenced model performance:

4.3.4.1 Positive Factors (High Accuracy Cases)

- **Well-defined binding pockets:** Proteins with deep, well-structured binding sites yielded more accurate predictions due to clearer structural features
- **Common ligand scaffolds:** Molecules with frequently occurring chemical motifs (e.g., aromatic rings, common functional groups) were predicted more accurately, likely due to better representation in the training data
- **Moderate binding affinity range:** Complexes with binding affinities in the middle range (pK_d 5–9) showed better prediction accuracy

- **High-resolution crystal structures:** Complexes derived from high-quality experimental structures (resolution $< 2.0 \text{ \AA}$) consistently showed lower prediction errors

4.3.4.2 Negative Factors (Reduced Accuracy Cases)

- **Flexible binding sites:** Proteins with large conformational flexibility or multiple binding modes led to higher prediction errors, as the static graph representation failed to capture dynamic behavior
- **Metal-containing complexes:** Ligands coordinating metal ions (e.g., Zn^{2+} , Mg^{2+}) were poorly predicted, likely due to inadequate representation of coordination chemistry in the current feature set
- **Large ligands:** Molecules exceeding 50 heavy atoms showed decreased accuracy, possibly due to increased graph complexity and higher-order interactions not fully captured by the model
- **Rare protein families:** Proteins from underrepresented families in the training set (e.g., certain kinases, viral proteases) exhibited lower prediction accuracy, indicating limited generalization
- **Water-mediated interactions:** Since water molecules were removed during pre-processing, complexes where water bridges play a critical role in binding showed reduced accuracy
- **Allosteric binding sites:** Non-orthosteric binding pockets with unconventional interaction patterns were challenging for the model to learn

4.4 Analysis and Discussion

This section provides an in-depth analysis of the experimental results, discusses the implications of the findings, and identifies challenges encountered during the initial implementation of ProtAI.

4.4.1 Performance Analysis

The initial results demonstrate that ProtAI achieves competitive performance compared to most baseline methods, with an RMSE of 1.55 kcal/mol and a Pearson correlation of 0.735. While the model performs slightly below the GraphSAGE baseline, these results

are still encouraging for an initial implementation and suggest that the graph neural network approach effectively captures key protein–ligand interaction patterns. The performance gap highlights specific areas for architectural refinement and feature engineering improvements.

4.4.1.1 Comparison with Baseline Methods

ProtAI outperformed traditional docking methods by a significant margin (29% reduction in RMSE compared to AutoDock Vina), validating the hypothesis that learned representations can better capture binding affinity than physics-based scoring functions. The improvement over simpler machine learning approaches (Random Forest, 18% RMSE reduction) and basic GCN architectures (5% RMSE reduction) demonstrates progress, though the model falls slightly short of the GraphSAGE baseline by approximately 5% in RMSE.

The Pearson correlation of 0.735 indicates that ProtAI reasonably ranks compounds by binding strength, which is valuable for virtual screening applications where relative ordering is important. However, the lower performance compared to GraphSAGE suggests that current architectural choices may not fully exploit the available structural information, particularly in modeling long-range dependencies and handling diverse binding modes.

4.4.1.2 Limitations and Error Analysis

Despite competitive overall performance, error analysis revealed systematic limitations:

1. **Static structure limitation:** The current implementation uses static protein structures, failing to account for conformational changes upon ligand binding (induced fit). This is a fundamental limitation that affects accuracy for flexible proteins and may explain both the lower performance on certain protein families and the gap with GraphSAGE, which may better handle structural variability through its aggregation scheme.
2. **Feature representation gaps:** The current atom-level features do not adequately capture quantum mechanical effects such as polarization, charge transfer, and dispersion interactions. Metal coordination chemistry is also poorly represented, leading to systematic errors for metalloproteins. This limitation particularly affects prediction accuracy for complexes with unusual electronic configurations.
3. **Message-passing depth limitations:** The current 3-layer GNN architecture may be insufficient for capturing long-range interactions in large binding pockets. GraphSAGE’s superior performance suggests that deeper or more sophisticated aggrega-

tion strategies could improve ProtAI’s ability to integrate information from distant atomic neighborhoods.

4. **Data imbalance:** The training dataset is dominated by certain protein families (e.g., kinases, proteases), leading to biased learning. Underrepresented proteins show higher prediction errors, indicating the need for more balanced training data or specialized augmentation techniques.
5. **Solvent effects:** Removing water molecules during preprocessing eliminates critical information about water-mediated hydrogen bonds and entropic contributions from water displacement. This simplification reduces accuracy for complexes where structured water molecules play key roles.

4.4.2 Challenges Encountered

Several technical and methodological challenges were identified during implementation:

4.4.2.1 Data Processing Challenges

- **Structure quality inconsistencies:** Many protein structures in the dataset contained missing residues, alternate conformations, and crystallographic artifacts that required manual curation or automated filtering. Approximately 8% of structures were excluded due to quality issues.
- **Ligand protonation states:** Determining correct protonation states at physiological pH proved challenging. Different protonation states can significantly affect binding affinity, but current preprocessing pipelines rely on heuristic rules that may not always be accurate.
- **Graph construction complexity:** Defining appropriate edge connections in the heterogeneous graph (protein–ligand interactions) required careful threshold tuning. Too restrictive thresholds missed important long-range interactions, while too permissive thresholds created overly dense graphs that slowed training.

4.4.2.2 Model Training Challenges

- **Memory limitations:** Large protein–ligand complexes (>5,000 atoms) occasionally exceeded GPU memory capacity during batch processing, requiring dynamic batching strategies or gradient accumulation techniques.

- **Overfitting tendency:** Initial experiments showed overfitting on the training set, particularly for rare protein families. This was partially mitigated through dropout, weight decay, and early stopping, but further regularization may be needed.
- **Hyperparameter sensitivity:** Model performance showed significant sensitivity to hyperparameters such as learning rate, number of graph convolution layers, and attention head count. Systematic hyperparameter optimization was time-intensive.
- **Training instability:** Some training runs exhibited unstable gradients or loss spikes, particularly in early epochs. Gradient clipping and learning rate warm-up strategies were implemented to address this issue.

4.4.2.3 Evaluation Challenges

- **Ground truth quality:** Experimental binding affinity measurements from different sources showed inconsistencies due to varying assay conditions, making it difficult to establish definitive ground truth values for some complexes.
- **Cross-dataset generalization:** Preliminary cross-dataset evaluation (training on PDBbind, testing on BindingDB) revealed performance degradation, indicating that the model may have learned dataset-specific biases rather than generalizable interaction patterns.
- **Interpretability limitations:** While attention weights provide some insight into important atomic interactions, extracting clear biological interpretations remains challenging. Further visualization and analysis tools are needed.

4.4.3 Future Improvements

Based on the analysis of results and challenges, several improvements are planned for future iterations:

1. **Architectural enhancements:** Increase message-passing depth to 4–5 layers and implement residual connections to capture long-range dependencies more effectively. Investigate GraphSAGE-style neighborhood sampling strategies to improve feature aggregation.
2. **Incorporate dynamic information:** Integrate molecular dynamics-derived features or ensemble structures to better capture protein flexibility and conformational changes, addressing the static structure limitation.

3. **Enhanced feature engineering:** Add quantum mechanical descriptors (e.g., partial charges from QM calculations, electrostatic potential surfaces) and explicit metal coordination features to improve predictions for metalloproteins and unusual binding modes.
4. **Advanced pooling mechanisms:** Implement hierarchical or attention-weighted pooling strategies to better aggregate atomic information into molecular-level representations, potentially closing the performance gap with GraphSAGE.
5. **Data augmentation strategies:** Implement structure-aware augmentation techniques such as random rotations, small perturbations to atomic coordinates, and synthetic complex generation to balance the dataset and improve generalization.
6. **Multi-task learning:** Extend the model to simultaneously predict binding affinity, binding pose quality, and interaction types, which may improve generalization through shared representations.
7. **Attention mechanism refinement:** Develop more sophisticated attention mechanisms that can explicitly identify and weight specific interaction types (hydrogen bonds, hydrophobic contacts, -stacking).
8. **Cross-dataset validation:** Conduct comprehensive cross-dataset experiments to assess generalization capability and identify dataset-specific biases that need to be addressed.

4.4.4 Summary

The initial experiments with ProtAI demonstrate promising results, achieving competitive performance on binding affinity prediction with an RMSE of 1.55 kcal/mol and a Pearson correlation of 0.735. The model successfully outperforms traditional docking methods (29% RMSE improvement), Random Forest baselines (18% RMSE improvement), and basic GCN architectures (5% RMSE improvement), validating the graph neural network approach.

However, ProtAI’s performance falls slightly short of the GraphSAGE baseline by approximately 5% in RMSE, indicating room for improvement. Several challenges remain, including limitations in representing protein flexibility, insufficient message-passing depth for long-range interactions, handling rare protein families, and capturing quantum mechanical effects. The performance gap with GraphSAGE, combined with detailed error analysis, provides a clear roadmap for future development, with specific areas requiring attention such as architectural deepening, dynamic structure integration, enhanced feature engineering, and improved data balancing.

These initial results establish a solid foundation for the continued development of ProtAI. While the model does not yet achieve state-of-the-art performance, the insights gained from error analysis and the identified architectural limitations provide concrete directions for improvement in subsequent phases of the project. The systematic evaluation against multiple baselines offers valuable understanding of where ProtAI excels and where further refinement is needed.

Chapter 5

Conclusions and Future Work

This chapter summarizes the key findings of the ProtAI research project, highlights the contributions made during the initial development phase, and outlines future directions for enhancing the framework’s performance and applicability in drug discovery.

5.1 Summary of Research

This research addressed the critical challenge of protein–ligand binding affinity prediction by developing **ProtAI**, a graph neural network-based framework designed to accelerate early-stage drug discovery. Traditional computational methods such as molecular docking rely on static protein structures and physics-based scoring functions that fail to capture the dynamic and complex nature of molecular interactions. Recent advances in deep learning, particularly graph neural networks, offer a promising alternative by learning directly from molecular structures and experimental binding data.

ProtAI was developed with the goal of predicting binding affinities more accurately than traditional methods while maintaining computational efficiency. The framework processes protein–ligand complexes as graph-structured data, where atoms serve as nodes and bonds represent edges, enabling the model to learn spatial and chemical interaction patterns through message-passing neural networks.

5.2 Key Findings

The initial implementation and evaluation of ProtAI yielded several important findings:

1. **Competitive performance over traditional methods:** ProtAI achieved an RMSE of 1.55 kcal/mol and a Pearson correlation coefficient of 0.735 on the PDBbind

test set, demonstrating a 29% improvement in RMSE compared to traditional docking methods (AutoDock Vina) and substantial gains over simpler machine learning baselines such as Random Forest (18% improvement) and basic GCN architectures (5% improvement).

2. **Validation of graph-based learning approach:** The results confirm that graph neural networks can effectively learn protein–ligand interaction patterns from structural data, providing a data-driven alternative to physics-based scoring functions. The model successfully captures local atomic environments and pairwise interactions critical to binding affinity.
3. **Performance gap with advanced baselines:** While ProtAI demonstrates progress, it performs slightly below the GraphSAGE baseline (5% higher RMSE), indicating that current architectural choices—particularly message-passing depth and aggregation strategies—require further refinement to achieve state-of-the-art performance.
4. **Identified systematic limitations:** Comprehensive error analysis revealed specific challenges, including difficulties in handling flexible binding sites, metal-containing complexes, large ligands, and underrepresented protein families. These limitations stem from the use of static structures, insufficient feature representations for quantum mechanical effects, and data imbalance.
5. **Importance of data quality and preprocessing:** Structure quality, protonation state determination, and graph construction strategies significantly impact model performance. Approximately 8% of structures required exclusion due to quality issues, highlighting the need for robust preprocessing pipelines.

5.3 Research Contributions

This research makes several contributions to the field of AI-driven drug discovery:

5.3.1 Technical Contributions

- **Implementation of end-to-end GNN pipeline:** Developed a complete pipeline from molecular data preprocessing (using RDKit, Biopython, and ProDy) to graph representation learning (using PyTorch Geometric) and model evaluation. This pipeline is modular and extensible, allowing for future architectural modifications and dataset integration.

- **Comprehensive baseline comparison:** Established systematic benchmarking against multiple baseline methods, including traditional docking, classical machine learning, and advanced GNN architectures, providing clear context for ProtAI’s performance and identifying areas for improvement.
- **Detailed error analysis framework:** Conducted thorough investigation of prediction failures across different molecular and structural categories, revealing insights into factors that enhance or reduce prediction accuracy. This analysis provides actionable guidance for future model refinements.
- **Reproducible experimental setup:** Documented hardware configurations, software environments, hyperparameters, and evaluation metrics to ensure reproducibility. Training scripts and preprocessing pipelines are organized for easy adaptation and extension.

5.3.2 Methodological Insights

- **Validation of graph-based molecular representation:** Demonstrated that graph neural networks can learn meaningful representations of protein–ligand complexes, supporting the growing body of evidence that deep learning approaches can complement or replace traditional computational chemistry methods.
- **Identification of architectural limitations:** Analysis revealed that shallow message-passing architectures (3 layers) may be insufficient for capturing long-range dependencies in large binding pockets, suggesting the need for deeper networks or hierarchical aggregation strategies.
- **Understanding of dataset biases:** Identified protein family imbalances in training data that lead to biased predictions, emphasizing the importance of balanced dataset curation and augmentation strategies for generalization.

5.4 Limitations and Challenges

Despite encouraging initial results, several limitations were encountered during the development of ProtAI:

1. **Suboptimal performance compared to state-of-the-art:** ProtAI’s performance lags behind more sophisticated GNN architectures such as GraphSAGE, indicating that further architectural innovations are necessary to close this gap.

2. **Static structure representation:** The current approach relies on single static conformations, failing to capture protein flexibility and conformational changes that occur during ligand binding (induced fit), which is a known limitation for accurate affinity prediction.
3. **Limited feature expressiveness:** Atom-level features do not adequately represent quantum mechanical effects such as polarization, charge transfer, and metal coordination chemistry, leading to systematic errors for certain complex types.
4. **Computational constraints:** Large protein–ligand complexes occasionally exceeded GPU memory limits, requiring dynamic batching and gradient accumulation techniques that increased training time.
5. **Generalization concerns:** Preliminary cross-dataset experiments (not fully reported) revealed performance degradation when training on PDBbind and testing on BindingDB, suggesting the model may have learned dataset-specific biases rather than generalizable interaction patterns.
6. **Interpretability challenges:** While the model provides attention weights, extracting clear biological interpretations of which specific interactions drive predictions remains difficult, limiting the model’s utility for mechanistic insights.

5.5 Future Work

Based on the findings and limitations identified in this initial research phase, several directions for future development have been outlined:

5.5.1 Short-Term Improvements (Next 3–6 Months)

1. **Architectural enhancements:**
 - Increase message-passing depth to 4–5 layers with residual connections to enable learning of long-range dependencies without gradient vanishing issues.
 - Implement GraphSAGE-style neighborhood sampling strategies to improve feature aggregation and computational efficiency.
 - Introduce hierarchical or attention-weighted global pooling mechanisms to better aggregate atomic information into molecular-level representations.
2. **Advanced feature engineering:**

- Integrate quantum mechanical descriptors such as partial charges from semi-empirical QM calculations, electrostatic potential surfaces, and frontier molecular orbital properties.
- Develop explicit representations for metal coordination chemistry, including coordination geometry and electron donation patterns.
- Incorporate physicochemical descriptors for binding pockets, such as hydrophobicity, electrostatic potential, and shape complementarity.

3. Data augmentation and balancing:

- Implement structure-aware augmentation techniques including random rotations, small coordinate perturbations, and conformational sampling to increase dataset diversity.
- Develop oversampling or class-weighting strategies to address protein family imbalances in the training set.
- Expand the dataset with additional sources such as BindingDB and ChEMBL to improve coverage of chemical and protein space.

5.5.2 Medium-Term Research Directions (6–12 Months)

1. Incorporation of dynamic information:

- Integrate molecular dynamics (MD) simulation-derived features such as RMSD, RMSF, and hydrogen bond occupancy to capture conformational flexibility.
- Explore ensemble-based approaches that train on multiple conformations per complex to account for structural variability.
- Investigate graph-based representations of dynamic trajectories, treating temporal evolution as a sequence learning problem.

2. Multi-task learning framework:

- Extend ProtAI to simultaneously predict binding affinity, binding pose quality scores, and interaction type classifications (hydrogen bonds, hydrophobic contacts, π -stacking).
- Leverage auxiliary tasks to learn more robust and generalizable molecular representations through shared encoder layers.
- Investigate whether multi-task learning can improve performance on under-represented protein families through knowledge transfer.

3. Attention mechanism refinement:

- Develop multi-head attention layers that explicitly model different interaction types (electrostatic, van der Waals, hydrogen bonding).
- Implement cross-attention mechanisms between protein and ligand subgraphs to better capture inter-molecular dependencies.
- Enhance interpretability by visualizing attention weights and correlating them with known binding interactions from experimental structures.

4. Comprehensive cross-dataset validation:

- Conduct systematic evaluations on external test sets (BindingDB, CASF, CSAR) to assess generalization capability.
- Perform domain adaptation experiments to understand and mitigate dataset-specific biases.
- Analyze performance across different protein families, ligand sizes, and binding affinity ranges to identify systematic failure modes.

5.5.3 Long-Term Vision (Beyond 12 Months)

1. Integration with virtual screening pipelines:

- Deploy ProtAI as a rescoring module in large-scale virtual screening campaigns to prioritize compounds for experimental validation.
- Optimize inference speed and memory efficiency to enable processing of millions of candidate compounds.
- Develop confidence estimation mechanisms to flag predictions with high uncertainty for manual review.

2. Explainable AI capabilities:

- Implement gradient-based attribution methods (e.g., integrated gradients) to identify critical atoms and interactions driving predictions.
- Develop visualization tools that overlay predicted interaction strengths on 3D molecular structures for intuitive interpretation by medicinal chemists.
- Validate model explanations against experimental mutagenesis data and structure-activity relationship (SAR) studies.

3. Extension to related prediction tasks:

- Adapt ProtAI for protein-protein interaction prediction, a related but distinct challenge in structural biology.

- Explore applications in predicting drug resistance mutations by analyzing how structural changes affect binding affinity.
- Investigate generative modeling approaches to propose novel ligand modifications that optimize predicted binding affinity.

4. Community resource development:

- Release ProtAI as an open-source tool with comprehensive documentation, pretrained models, and example workflows.
- Develop a web-based interface for non-expert users to submit protein–ligand complexes and receive binding affinity predictions.
- Contribute datasets, benchmarking protocols, and evaluation metrics to the computational drug discovery community.

5.6 Concluding Remarks

This research represents the initial phase of developing ProtAI, a graph neural network framework for protein–ligand binding affinity prediction. While the model demonstrates competitive performance over traditional methods and validates the graph-based learning approach, the results also reveal clear limitations and opportunities for improvement. ProtAI currently achieves an RMSE of 1.55 kcal/mol, which surpasses traditional docking methods but falls slightly short of more advanced GNN baselines.

The systematic evaluation, comprehensive error analysis, and identification of architectural limitations provide a strong foundation for future development. The performance gap with state-of-the-art models, rather than being discouraging, offers concrete directions for refinement, including deeper message-passing architectures, richer feature representations, and incorporation of dynamic structural information.

Looking forward, the planned improvements—spanning architectural enhancements, advanced feature engineering, multi-task learning, and integration of molecular dynamics—have the potential to significantly boost ProtAI’s predictive accuracy and generalization capability. Beyond performance metrics, enhancing interpretability through attention visualization and gradient-based attribution will be crucial for ProtAI’s adoption in real-world drug discovery workflows, where understanding *why* a model makes certain predictions is as important as the predictions themselves.

Ultimately, this research contributes to the growing body of work demonstrating that deep learning approaches can learn complex biochemical interaction patterns from data, offering a promising complement to traditional computational chemistry methods. The insights gained from this initial implementation—both successes and limitations—will

guide the continued development of ProtAI toward becoming a practical, reliable tool for accelerating the early stages of drug discovery and reducing the time and cost associated with identifying promising therapeutic candidates.